

Long Page Chunking Test

This page is intentionally designed to be very long so that the PDF parser must split it into multiple chunks of 512 characters each.

The purpose of this test is to verify that the chunking logic inside `chunk_text` works correctly and that the parser does not assume that a single page always equals a single chunk.

We need to ensure that `global_pdf_chunk_index` increases continuously across chunks and that no chunk exceeds the maximum allowed size.

Furthermore, this test helps confirm that metadata such as `page_number` and chunk ordering remain consistent even when a single page produces many chunks.

The system should correctly handle this without errors and without losing any content during splitting.