

Page 1 - Introduction

This is the first page of a multi-page PDF used for testing.

It contains structured text to validate that page separation works correctly.

Each page should be processed independently by the parser.

Page 2 - Content Processing

This page continues the test document.

It ensures that chunking works correctly across multiple pages.

The parser should reset page context but continue global indexing.

Page 3 - Final Section

This is the last page of the test PDF.

It confirms that the parser handles end-of-document cases properly.

No additional pages should be present after this one.